
SUMMARY

Data Engineer with 4+ years building scalable pipelines and cloud-native architectures. Migrated 10TB+ datasets, optimized performance by 60%, and engineered fault-tolerant ETL/ELT systems processing millions of records using Python, SQL, Snowflake, AWS, and Airflow.

SKILLS & CERTIFICATIONS

- **Programming Languages:** Python, SQL, Shell Scripting
- **Cloud & Data Platforms:** AWS (S3, Lambda, Glue, EC2), Snowflake, BigQuery, Redshift, Oracle, GCP, Databricks
- **Tools & DevOps:** Git, Apache Spark, Apache Airflow, dbt, Pandas, Iceberg, JIRA, Confluence, VS Code
- **Practices:** Agile/Scrum, CI/CD, Code Reviews, Object-Oriented Design
- **Certifications:** AWS Certified Data Engineer - Associate, SnowPro Core Certified

PROFESSIONAL EXPERIENCE

Data Engineer Capgemini

May 2025 – Present
Atlanta, Georgia

- Designed centralized error-handling framework in Snowflake using Snowpark and Python, reducing pipeline failures by 30% and improving observability through granular logging and monitoring.
- Redesigned ingestion framework into modular, metadata-driven distributed pipeline leveraging Python and SQL, improving runtime efficiency for S3-to-Snowflake ETL workflows supporting 800+ data sources.
- Built secure ingestion pipelines for SFTP sources using Python that orchestrate transformations from S3 to Snowflake, processing 1M+ daily records and reducing data source onboarding time by 50% through standardized configuration and comprehensive testing.
- Integrated 20+ external REST APIs using Python into data exchange framework, implementing retry logic, rate limiting, and error handling that achieved 99.5% data retrieval success rate.
- Developed dbt models following software engineering best practices including Git version control, code reviews, and CI/CD testing, enabling analytics for 100+ stakeholders.

Data Engineer Intern

January 2025 – April 2025
Austin, Texas

Texas Department of Motor Vehicles

- Streamlined daily ingestion of 500K+ records from DB2 using Python and SQL, replacing manual processes and reducing data delivery time by 80%.
- Architected Informatica workflow to migrate 22M+ records from DB2 to SQL Server with data validation and error handling, improving reporting performance

Data Engineer Capgemini

June 2021 – June 2023
Mumbai, Maharashtra

- Orchestrated end-to-end migration of 10TB to Snowflake cloud platform, designing scalable ETL architecture that improved query response time by 60% for 100+ users with zero data loss.
- Engineered an ELT pipeline with Python and Airflow, streamlining data loading processes and reducing manual intervention by 30%.
- Programmed and optimized SQL queries, reducing processing time by 40% and enhancing data retrieval.
- Collaborated with cross-functional teams to define data requirements and optimize reporting structures, leading to a 40% improvement in overall reporting efficiency.

PROJECT EXPERIENCE

IMDb Data Warehouse: ELT Pipeline | Python, SQL, Data Modeling, dbt, BigQuery, REST APIs

- Built a fully normalized movies data warehouse from scratch by collecting and enriching datasets via IMDb APIs; cleaned and standardized data using Gemini and designed schema with strong referential integrity.
- Engineered a modular ELT pipeline in BigQuery using dbt with layered transformations (staging, intermediate, reporting); scheduled via shell scripts, and added change data capture with dbt snapshots for historical tracking.

EDUCATION

University of Texas at Austin

Master of Science in Information Systems Management | **GPA: 3.94/4.0**

May 2025
Austin, Texas

- Specializing in Cloud Computing & Data Engineering

Thakur College of Engineering and Technology, University of Mumbai

Bachelor of Engineering, Information Technology | **GPA: 3.3/4.0**

May 2021
Mumbai, Maharashtra